

Lecture 3b: Taxonomy of RL Algorithms, Q-Learning

Qi Zhang

Last Updated: October 2025

1 Taxonomy of RL Algorithms

In this note, we begin studying RL algorithms. We will focus on the infinite-horizon, discounted-reward setting where the MDP is specified by a tuple $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ with initial state distribution $d_0 \in \Delta(\mathcal{S})$. The RL agent, or learner, does not know (P, R) a priori and interacts with the MDP following Protocol 1, where the role of the learner's RL algorithm is highlighted in green.

Protocol 1 RL interaction (infinite-horizon)

```

1: learner initializes and will maintain its internal state;
2: repeat over episodes:
3:   learner observes initial state  $s_0 \sim d_0$  of the current episode;
4:   for  $h = 1, \dots, H$  until termination or truncation do    ▷ truncation:  $s_{H+1}$  is not terminal
5:     learner chooses action  $a_h$  based on its internal state;
6:     learner takes  $a_h$  and observes next state  $s_{h+1} \sim P(s_h, a_h)$  and reward  $r_h = R(s_h, a_h)$ ;
7:     learner updates its internal state;
8:   end for
9: until some stopping criterion is met
10: learner outputs a policy  $\hat{\pi}$  based on its internal state;
```

Crucially, the learner maintains its *internal state* that 1) gets updated based on the newest transition (line 7) and 2) determines the very next action to take (line 5). This way, the learner is fully adaptive: in general, it chooses the current action based on all previous transitions. The learner's internal state can include some or all of the following components:

- Value estimates: $V : \mathcal{S} \rightarrow \mathbb{R}, Q : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$
- Policy/actor: $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$
- Replay buffer: $\mathcal{D} = (s, a, r, s')$ storing previous transitions
- Model estimate: (\hat{P}, \hat{R}) that estimates the ground-truth (P, R) based on previous transitions

RL algorithms are often characterized by what components above are (not) included in the learner's internal state. Table 1 provides a (simplified) taxonomy, although it is usually of little importance to memorize such a table.

2 Tabular Q-Learning

We now introduce our first RL algorithm, *Q-Learning*. Its key idea is that, in order to find an optimal policy (by the end of Protocol 1 at line 10), it suffices to find out the optimal Q-value Q^* because an optimal policy can be then extracted by acting greedily with respect to it, i.e., $\pi^*(s) = \arg \max_{a \in \mathcal{A}} Q^*(s, a)$. In Q-learning, the learner's internal state includes only a Q-value

Table 1: A Taxonomy of RL Algorithms.

	Value estimates V, Q	Policy π	Model estimate (\hat{P}, \hat{R})
Model-free	-	-	Y
Model-based	-	-	Y
Value-based	Y	N	N
Policy-based	N	Y	N
Actor-critic	Y	Y	N
Tabular	Represented by tables, one entry for each (s) or (s, a)		
Function Approximation	Represented by generic functions		

estimate $Q : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ and therefore it is a value-based (and also model-free) RL algorithm. The goal of Q-learning is to update Q such that it eventually approximates Q^* well and recommends the outputs the policy as $\hat{\pi}(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$. *Tabular* Q-learning represents Q as a table/vector with $|\mathcal{S} \times \mathcal{A}|$ entries, one entry per (s, a) pair, and, upon a transition (s_t, a_t, r_t, s_{t+1}) , updates the entry of (s_t, a_t) as

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha_t \left(r_t + \gamma \max_{a \in \mathcal{A}} Q(s_{t+1}, a) - Q(s_t, a_t) \right). \quad (1)$$

Here, the index t is the total number of transitions the learner has experienced in Protocol 1, which keeps increasing across episodes; $\alpha_t \in (0, 1)$ controls how aggressive the update is, which in general depends on t .

To understand why update (1) makes sense, recall Q-value iteration: starting with an arbitrary $Q \in \mathbb{R}^{|\mathcal{S} \times \mathcal{A}|}$, update it iteratively as $Q \leftarrow \mathcal{T}Q$ where *Bellman optimality operator* $\mathcal{T} : \mathbb{R}^{|\mathcal{S} \times \mathcal{A}|} \rightarrow \mathbb{R}^{|\mathcal{S} \times \mathcal{A}|}$ is defined as:

$$(\mathcal{T}Q)(s, a) := R(s, a) + \gamma \mathbb{E}_{s' \sim P(s, a)} \left[\max_{a \in \mathcal{A}} Q(s', a) \right] \quad (2)$$

$$\approx r + \gamma \max_{a \in \mathcal{A}} Q(s', a) \quad (3)$$

where the \approx approximates the expectation by a single sample of transition (s, a, r, s') with $r = R(s, a)$ and $s' \sim P(s, a)$. In the planning setting, we can perform the Q-value iteration and we have the convergence of $\mathcal{T}^k Q \rightarrow Q^*$ as $k \in \infty$, where \mathcal{T}^k iteratively applies \mathcal{T} for k times. In the RL setting, we cannot perform a full value iteration update $(\mathcal{T}Q)$ because (P, R) is unknown. Instead, upon on a new transition (s, a, r, s') , we update $Q(s, a)$ to be closer to the one-sample target (3):

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \left(r + \gamma \max_{a \in \mathcal{A}} Q(s', a) \right) = Q(s, a) + \alpha \left(r + \gamma \max_{a \in \mathcal{A}} Q(s', a) - Q(s, a) \right)$$

where $\alpha \in (0, 1)$ controls how aggressively $Q(s, a)$ is updated to target (3). Setting $(s, a, r, s') = (s_t, a_t, r_t, s_{t+1})$ and $\alpha = \alpha_t$ in the update above recovers (1).

It turns out that, if the following conditions hold, we have the convergence guarantee of $Q \rightarrow Q^*$ as $t \rightarrow \infty$:

- Every (s, a) pair is visited infinitely often.

- α_t decreases but not too quickly.

A formal description of the first condition is $\lim_{t \rightarrow \infty} N_t(s, a) = \infty$ for any (s, a) , where $N_t(s, a)$ is the number of times (s, a) is visited by t . This means the learner must do *exploration* when choosing actions at line 5 in Protocol 1. For example, one can do ϵ -greedy w.r.t. the current Q . A formal description of the second condition is $\sum_t \alpha_t = \infty, \sum_t \alpha_t^2 < \infty$ (e.g., $\alpha_t = \frac{1}{t}$ satisfies this).

3 Deep Q-Network

The obvious issue of tabular Q-learning is the difficulty of scaling to large state/action spaces. *Function approximation* solves this issue by representing the Q-value estimates with a generic function. The intuition is that, by updating the Q-value on a specific state-action pair, the changes made to the function will also influence other state-action pairs.

Protocol 2 DQN algorithm

- 1: learner initializes parameter θ of the Q network and sets replay buffer $\mathcal{D} = \{\}$ and target network parameter $\bar{\theta} = \theta$;
 - 2: **repeat** over episodes:
 - 3: learner observes initial state $s \sim d_0$ of the current episode;
 - 4: **for** timesteps within the episode **do**
 - 5: learner chooses action a that is ϵ -greedy w.r.t. $Q_\theta(s, \cdot)$; $\triangleright \epsilon$ usually decreases over time
 - 6: learner takes a and observes next state s' and reward r ;
 - 7: learner add (s, a, r, s') to replay buffer \mathcal{D} ;
 - 8: learner samples a batch of N transitions from \mathcal{D} ;
 - 9: learner takes a gradient step minimizing $L(\theta)$;
 - 10: learner sets $\bar{\theta} \leftarrow \theta$ every c updates of θ ; \triangleright updates target network periodically
 - 11: **end for**
 - 12: **until** some stopping criterion is met
 - 13: learner outputs a policy $\hat{\pi}$ as the greedy policy w.r.t. Q_θ ;
-

A representative work is *Deep Q-Network* (DQN), where the Q-value estimates is represented using a neural network, $Q_\theta : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ where θ is the neural network's parameters. Although people tried neural networks for Q-learning before DQN, those prior efforts updated the parameters on the most recent transition, in a similar fashion as in tabular Q-learning, and observed limited success due to DQN improved the training (i.e., parameters updating) of the networks with the following innovations:

- It stores the past transitions in a memory, often referred to as *replay buffer*, $\mathcal{D} = \{(s, a, r, s')\}$.
- On a given transition (s, a, r, s') , it forms target (3) using another network $Q_{\bar{\theta}}$, which is of the same structure as Q_θ but with its parameter $\bar{\theta}$ updated more slowly than θ .
- To update θ , it samples a batch of transitions from the replay buffer and update θ with a single gradient step to reduce the Q-value estimates with the formed target values.

This leads to the following loss function for network Q_θ :

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \left(Q_\theta(s_i, a_i) - (r_i + \gamma \max_{a \in \mathcal{A}} Q_{\bar{\theta}}(s'_i, a)) \right)^2 \quad \text{where } (s_i, a_i, r_i, s'_i) \sim \mathcal{D} \text{ for } i = 1, \dots, N.$$